

[Skip to main content](#)

> [cs](#) > arXiv:2602.19919

 Search...

quick links

- [Login](#)
- [Help Pages](#)
- [About](#)

Computer Science > Computation and Language

arXiv:2602.19919 (cs)

[Submitted on 23 Feb 2026]

Janus-Q: End-to-End Event-Driven Trading via Hierarchical-Gated Reward Modeling

[Xiang Li](#), [Zikai Wei](#), [Yiyan Qi](#), [Wanyun Zhou](#), [Xiang Liu](#), [Penglei Sun](#), [Yongqi Zhang](#), [Xiaowen Chu](#)

[View PDF](#) [HTML \(experimental\)](#)

Financial market movements are often driven by discrete financial events conveyed through news, whose impacts are heterogeneous, abrupt, and difficult to capture under purely numerical prediction objectives. These limitations have motivated growing interest in using textual information as the primary source of trading signals in learning-based systems. Two key challenges hinder existing approaches: (1) the absence of large-scale, event-centric datasets that jointly model news semantics and statistically grounded market reactions, and (2) the misalignment between language model reasoning and financially valid trading behavior under dynamic market conditions. To address these challenges, we propose Janus-Q, an end-to-end event-driven trading framework that elevates financial news events from

auxiliary signals to primary decision units. Janus-Q unifies event-centric data construction and model optimization under a two-stage paradigm. Stage I focuses on event-centric data construction, building a large-scale financial news event dataset comprising 62,400 articles annotated with 10 fine-grained event types, associated stocks, sentiment labels, and event-driven cumulative abnormal return (CAR). Stage II performs decision-oriented fine-tuning, combining supervised learning with reinforcement learning guided by a Hierarchical Gated Reward Model (HGRM), which explicitly captures trade-offs among multiple trading objectives. Extensive experiments demonstrate that Janus-Q achieves more consistent, interpretable, and profitable trading decisions than market indices and LLM baselines, improving the Sharpe Ratio by up to 102.0% while increasing direction accuracy by over 17.5% compared to the strongest competing strategies.

Subjects: **Computation and Language (cs.CL)**; Machine Learning (cs.LG)

Cite as: [arXiv:2602.19919](#) [cs.CL]

(or [arXiv:2602.19919v1](#) [cs.CL] for this version)

<https://doi.org/10.48550/arXiv.2602.19919>

 Focus to learn more

arXiv-issued DOI via DataCite (pending registration)

Submission history

From: Xiang Li [[view email](#)]

[v1] Mon, 23 Feb 2026 14:58:51 UTC (2,984 KB)

Bibliographic Tools

Bibliographic and Citation Tools

Bibliographic Explorer Toggle

Bibliographic Explorer ([What is the Explorer?](#))

Connected Papers Toggle

Connected Papers ([What is Connected Papers?](#))

Litmaps Toggle

Litmaps ([What is Litmaps?](#))

scite.ai Toggle

scite Smart Citations ([What are Smart Citations?](#))

Code, Data, Media

Code, Data and Media Associated with this Article

alphaXiv Toggle

alphaXiv ([What is alphaXiv?](#))

Links to Code Toggle

CatalyzeX Code Finder for Papers ([What is CatalyzeX?](#))

- DagsHub Toggle
DagsHub ([What is DagsHub?](#))
- GotitPub Toggle
Gotit.pub ([What is GotitPub?](#))
- Huggingface Toggle
Hugging Face ([What is Huggingface?](#))
- Links to Code Toggle
Papers with Code ([What is Papers with Code?](#))
- ScienceCast Toggle
ScienceCast ([What is ScienceCast?](#))
- Demos

Demos

- Replicate Toggle
Replicate ([What is Replicate?](#))
- Spaces Toggle
Hugging Face Spaces ([What is Spaces?](#))
- Spaces Toggle
TXYZ.AI ([What is TXYZ.AI?](#))
- Related Papers

Recommenders and Search Tools

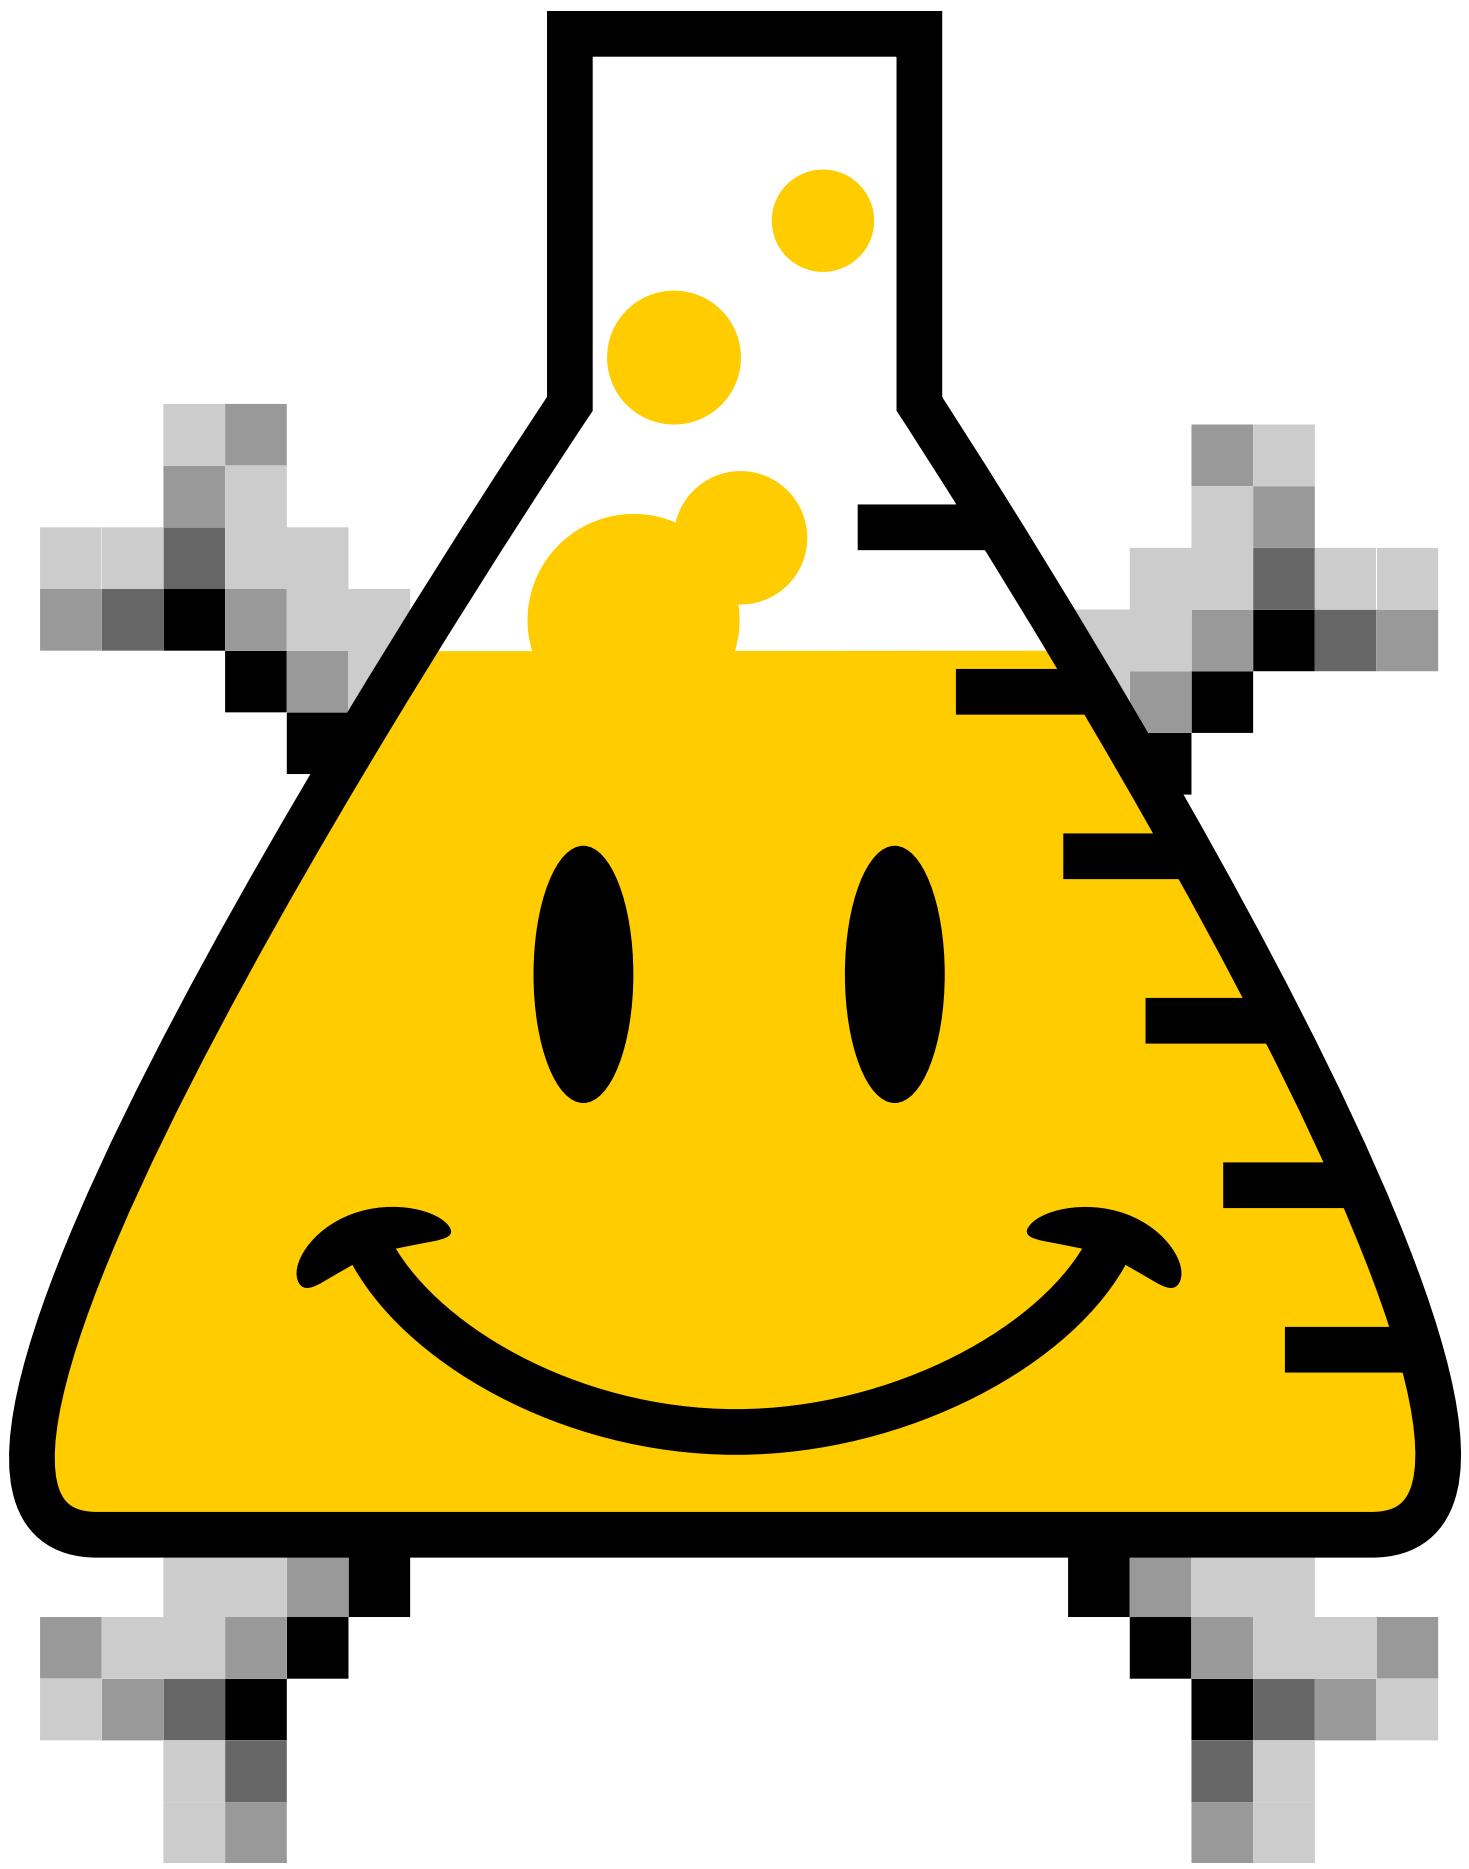
- Link to Influence Flower
Influence Flower ([What are Influence Flowers?](#))
- Core recommender toggle
CORE Recommender ([What is CORE?](#))
- About arXivLabs

arXivLabs: experimental projects with community collaborators

arXivLabs is a framework that allows collaborators to develop and share new arXiv features directly on our website.

Both individuals and organizations that work with arXivLabs have embraced and accepted our values of openness, community, excellence, and user data privacy. arXiv is committed to these values and only works with partners that adhere to them.

Have an idea for a project that will add value for arXiv's community? [Learn more about arXivLabs](#).



[Which authors of this paper are endorsers?](#) | [Disable MathJax](#) ([What is MathJax?](#))